Robust Active Contour Models for Noisy and Low-Contrast Image Segmentation

Effie Wu and Lyla Liu and Md Aeinul Islam and Jaehoon Choi
University of California, Los Angeles (UCLA)
{effiewu0628,xuanning16,mdaeinulislam,jaehoonchoi719}@g.ucla.edu

Abstract

Active contour models often struggle with noisy and low-contrast images because fixed parameters and single-scale optimization adapt poorly to changing boundaries. In this work, we improve contour robustness through three complementary strategies: spatially adaptive parameterization, coarse-to-fine multiscale optimization, and reinforcement-learning-based parameter scheduling. We evaluate Classical, Adaptive, Full-Adaptive, RL-Adaptive, Multi-Scale, and Combined snake variants on synthetic benchmarks, fluorescence microscopy, and ultrasound datasets. Results show that adaptive approaches consistently outperform the Classical Snake baseline. RL-Adaptive achieves the strongest overall performance on challenging synthetic and ultrasound images, while the Combined framework performs best on low-contrast fluorescence microscopy by improving contour fitting robustness. These findings demonstrate that integrating adaptive optimization strategies significantly improves active contour segmentation under difficult imaging conditions.

1 Introduction

Active contour models, commonly known as snakes, are widely used for image segmentation because they provide an interpretable mechanism for locating object boundaries. Although classical snakes perform well when boundaries are clear, their reliance on fixed parameters makes them sensitive to noise, low contrast, weak edges, and complex object geometry.

A single-parameter configuration is rarely optimal throughout contour evolution. Strong boundaries may require greater flexibility to capture local details, whereas noisy or weak-edge regions benefit from stronger regularization. Similarly, early optimization often requires aggressive movement toward the target, while later stages require conservative refinement. Furthermore, contour optimization is sensitive to initialization and may converge to local minima when boundaries are complex or the initial contour is far from the target. These spatial, temporal, and scale-related differences motivate active contour models that can adjust their behavior during optimization.

In this work, we investigate several adaptive extensions to classical snakes. Adaptive Snake varies bending stiffness according to local gradient magnitude, while Full-Adaptive Snake additionally adapts damping at each contour point. RL-Adaptive Snake combines spatially adaptive stiffness with a Deep Q-Network that dynamically schedules damping and smoothing scale parameters during contour evolution. We further examine multiscale optimization, in which contours evolve from coarse to fine image resolutions to avoid local minima, and a Combined framework that integrates coarse-to-fine optimization with spatially adaptive stiffness.

We evaluate these methods on five synthetic test cases and two biomedical datasets containing fluorescence microscopy and ultrasound images. Results show that RL-Adaptive improves IoU from 0.918 to 0.958 on the challenging noisy-star case and from 0.615 to 0.644 on ultrasound images. On fluorescence microscopy images, the Combined framework achieves the strongest performance, suggesting that coarse-to-fine initialization provides particular advantages for datasets with compact and consistent object structures. Overall, our results demonstrate that spatial adaptation, temporal scheduling, and multiscale optimization are complementary strategies whose relative benefits depend on image characteristics.

Our main contributions are:

- spatially adaptive snake formulations for contour stiffness and damping;
- an RL-based controller for temporal optimization-parameter scheduling;

- a Combined framework integrating coarse-to-fine multiscale optimization with spatial adaptation; and

- a comparative evaluation across synthetic and biomedical segmentation tasks.

2 Background

2.1 Classical Active Contour Models

Active contour models were first introduced by Kass et al. (1) as a deformable curve framework for image segmentation and boundary extraction. Starting from an initial contour, the snake evolves by minimizing an energy functional that balances internal smoothness constraints with image-driven external forces.

Given a parameterized contour $c(u)$, the snake energy is defined as

$$E(c) = E_{\text{int}}(c) + E_{\text{ext}}(c), \tag{1}$$

where E_{int} denotes the internal energy and E_{ext} denotes the external image energy.

The internal energy is typically expressed as

$$E_{\text{int}}(c) = \frac{1}{2}\int_0^1 \left[\alpha \left|\frac{\partial c}{\partial u}\right|^2 + \beta \left|\frac{\partial^2 c}{\partial u^2}\right|^2\right] du, \tag{2}$$

where α controls contour tension and β controls contour rigidity. Larger values of α encourage contour contraction, while larger values of β penalize curvature and promote smoother shapes.

The external energy attracts the contour toward image features such as edges and object boundaries. In gradient-based snakes, this energy is commonly derived from image gradient magnitude, causing the contour to converge toward regions with strong edge responses.

Although classical snakes provide an interpretable segmentation framework, the performance is generally sensitive to noise, weak boundaries, and parameter selection. These limitations motivate the proposed active and multiscale extensions explored in this work.

2.2 Dynamic Snake Formulation

The dynamic snake formulation introduces temporal evolution and damping into contour optimization.

The continuous formulation can be written as

$$D\frac{\partial c}{\partial t} + Kc = f \tag{3}$$

where K represents the internal force operator, D is the damping matrix, and f denotes external image forces.

After temporal discretization, the update equation becomes

$$\left(K + \frac{D}{\Delta t}\right) c^{t+1} = f^t + \frac{D}{\Delta t} c^t \tag{4}$$

In the classical formulation, the damping term is represented as

$$\frac{D}{\Delta t} = \gamma I \tag{5}$$

where γ is a constant damping coefficient and I is the identity matrix.

3 Method

3.1 Adaptive Energy Weighting

The classical snake model employs fixed internal energy parameters throughout contour evolution. While this formulation provides stable optimization behavior, it cannot adapt to the varying characteristics of different image regions. Strong boundaries often require increased contour flexibility to capture fine geometric details, whereas weak or noisy regions benefit from stronger regularization to suppress spurious deformations.

To address this limitation, we investigate a series of adaptive snake formulations that modify the internal model parameters according to local image structure. We begin with a spatially adaptive stiffness model, extend it to a full adaptive formulation with adaptive damping, and finally introduce a reinforcement-learning-based framework that performs temporal parameter scheduling during contour evolution.

3.1.1 Adaptive Snake

The classical snake employs a fixed bending stiffness parameter throughout contour evolution. However, different image regions often require different levels of contour flexibility.

To address this limitation, we introduce a spatially adaptive bending stiffness parameter

$$\beta(u) = \beta_{\min} + (\beta_{\max} - \beta_{\min}) \exp\left(-k|\nabla I(c(u))|\right) \tag{6}$$

Here, I denotes the image intensity function and $|\nabla I(c(u))|$ serves as a local measure of edge strength along the contour.

Near strong image boundaries, the gradient magnitude becomes large, causing $\beta(u)$ to decrease. This reduces the curvature penalty and allows the contour to deform more freely to capture fine boundary structures. Conversely, in weak-edge or noisy regions, $\beta(u)$ increases, producing smoother contour evolution and stronger regularization.

Unlike adaptive elasticity formulations, we keep α fixed throughout optimization and only adapt β, since β directly controls contour flexibility and local boundary fitting.

3.1.2 Full-Adaptive Snake

To further improve contour evolution, we introduce a spatially adaptive damping parameter on top of the standard adaptive snake.

$$\gamma(u) = \gamma_{\min}+(\gamma_{\max}-\gamma_{\min}) \exp\left(-k_\gamma|\nabla I(c(u))|\right) \tag{7}$$

The classical damping term

$$\frac{D}{\Delta t} = \gamma I \tag{8}$$

is replaced by

$$\frac{D}{\Delta t} = diag\left(\gamma(u)\right) \tag{9}$$

This allows each contour point to possess its own damping coefficient based on local image structure.

Near strong boundaries, $\gamma(u)$ decreases, reducing the influence of the previous contour state and allowing the contour to respond more aggressively to image forces. In weak-edge regions, $\gamma(u)$ increases, resulting in more conservative updates and improved numerical stability.

The Full-Adaptive Snake therefore combines adaptive contour flexibility through $\beta(u)$ with adaptive update behavior through $\gamma(u)$.

3.1.3 RL-Adaptive Snake

Although spatial adaptation improves contour behavior across different image regions, the optimization strategy remains fixed throughout contour evolution.

To introduce temporal adaptation, we formulate parameter selection as a reinforcement learning problem and employ a Deep Q-Network (DQN) controller (2).

The contour continues to use the adaptive stiffness formulation introduced previously, while the DQN selects global optimization parameters during contour evolution.

The observation space consists of three groups of features:

- Local gradient magnitude and direction sampled at 32 probe points along the contour;
- Contour statistics, including mean radius, radius standard deviation, and centroid offset;
- Optimization state information, including iteration progress, current γ, current σ, and mean contour edge strength.

Every 50 snake iterations, the agent selects one action corresponding to a parameter pair (γ, σ).

The available action presets are

$$(1,2), (2,3), (5,5), (10,6), (20,8) \tag{10}$$

Here, γ controls update aggressiveness while σ controls the image smoothing scale used during force computation.

The reward is defined as

$$r_t = IoU_t - IoU_{t-1} \tag{11}$$

where IoU denotes the Intersection-over-Union segmentation metric. The reward, therefore, measures the improvement in segmentation quality between consecutive decision steps.

Positive rewards correspond to improved contour fitting, while negative rewards indicate deterioration in segmentation quality.

Conceptually, adaptive $\beta(u)$ provides spatial adaptation by determining where the contour should deform, while the DQN controller provides temporal adaptation by determining when the contour should evolve aggressively or conservatively.

3.2 Multi-Scale Optimization

Classical snake optimization is often sensitive to initialization and may converge to poor local minima in noisy or low-contrast images. To improve optimization stability, we introduce a coarse-to-fine multi-scale strategy that guides contour evolution through progressively finer image resolutions.

Given an input image, we construct a pyramid of downsampled representations. Contour optimization is first performed at the coarsest resolution, where high-frequency noise is suppressed, and the global object structure is easier to capture. After convergence at one scale, the resulting contour is upsampled and used to initialize optimization at the next finer scale. This process is repeated until the original image resolution is reached.

Unlike the adaptive and RL-based variants, the Multi-Scale Snake keeps the classical contour parameters fixed and modifies only the optimization procedure. Its purpose is therefore to isolate the effect of coarse-to-fine refinement without introducing spatially varying stiffness or temporally varying control parameters.

By beginning at coarse resolutions and progressively refining the contour, the multi-scale framework improves robustness to noise and reduces the likelihood of local-minimum failure. This makes the optimization process more stable while preserving the interpretability of the classical active contour formulation.

3.3 Combined Framework

The Combined Snake unifies the multi-scale spatial pyramid with the point-by-point adaptive stiffness (β) model. This framework changes its mechanical flexibility near boundaries and alters its behavior depending on the resolution layer. We introduce two engineering updates to stabilize this process on real medical scans:

3.3.1 Quantile-Based Force Scaling

Classical active contours normalize external force fields by dividing by the absolute maximum gradient value found in the entire image. However, medical images often contain bright artifacts or noise patches. If the absolute maximum belongs to an irrelevant noise speckle, the denominator becomes artificially inflated, shrinking the true boundary forces near the target object to almost zero.

To resolve this blinding effect, we implement a smooth activation framework using the hyperbolic tangent ($\tanh$) function:

$$f_{\text{robust}} = \tanh(f) \tag{12}$$

The $\tanh$ function naturally scales any input force vector into a continuous bounded range between -1.0 and 1.0. This formulation completely squashes extreme outlier noise spikes while preserving the linear behavior of soft, weak edges near the true object boundaries.

3.3.2 Resolution-Aware Variable Damping

Moving the contour across a downsampled image means that a 1-pixel step covers a large spatial distance relative to the original image scale. If the model uses a heavy, fixed damping parameter (γ) at all resolution levels, the contour experiences too much friction and cannot travel the long distance to the boundary.

We implement a dynamic damping parameter γ_m that automatically lightens the movement brake at lower resolutions:

$$\gamma_m = \max\left(\gamma_{\min},\, \gamma \cdot \omega^{L-1-m}\right) \tag{13}$$

Friction is reduced during the coarse stages to allow rapid travel from the initialization zone, while full precision damping is applied at the original resolution layer to lock the contour perfectly onto the fine edge.

4 Experimental Setup

4.1 Synthetic Benchmark Dataset

To evaluate the robustness of different snake formulations under controlled conditions, we constructed a synthetic benchmark dataset consisting of geometric shapes with varying levels of complexity.

The dataset includes circles, ellipses, and non-convex star-shaped objects. For each shape category, both clean and noisy variants were generated. The clean images evaluate the ability of each method to recover object boundaries under ideal conditions, while the noisy images assess robustness to image perturbations and weak edge information.

The benchmark was designed to progressively increase segmentation difficulty. Circular and elliptical objects primarily test basic contour fitting performance, whereas star-shaped objects contain sharp corners and concave regions that present greater challenges for active contour models. Ground-truth segmentation masks were generated directly from the synthetic shapes and were used for quantitative evaluation.

4.2 Real-World Datasets

To further evaluate the generalization ability of the proposed methods, we additionally consider two real-world biomedical image segmentation datasets.

4.2.1 FLUO Cell Segmentation Dataset

The FLUO Cell Segmentation Dataset from the Cell Tracking Challenge (Fluo-N2DL-HeLa)" (3) consists of fluorescence microscopy images with manually annotated cell boundaries. Compared with synthetic shapes, cell images exhibit greater geometric variability, irregular boundaries, and imaging noise, making accurate contour evolution significantly more challenging.

4.2.2 Kaggle Ultrasound Dataset

The Kaggle Ultrasound Dataset (4) contains ultrasound images characterized by low contrast, weak object boundaries, and speckle noise. These characteristics often degrade the performance of classical active contour models and therefore provide a challenging benchmark for evaluating the robustness of adaptive and reinforcement-learning-based contour optimization strategies.

4.3 Evaluation Metrics

We evaluate segmentation performance using three complementary metrics: Intersection-over-Union (IoU), Hausdorff Distance (HD), and Mean Boundary Distance (MBD).

IoU measures the overlap between the predicted segmentation mask and the ground-truth mask.

Hausdorff Distance evaluates the maximum discrepancy between the predicted contour and the ground-truth boundary, providing a measure of worst-case boundary error.

Mean Boundary Distance computes the average distance between contour points and the ground-truth boundary, reflecting overall contour localization accuracy.

For IoU, higher values indicate better segmentation performance, whereas lower Hausdorff Distance and Mean Boundary Distance values indicate more accurate boundary fitting.

4.4 Compared Methods

We compare six active contour variants:

- Classical Snake
- Adaptive Snake
- Full-Adaptive Snake
- RL-Adaptive Snake
- Multi-Scale Snake
- Combined Snake

The Classical Snake serves as the baseline method. Adaptive and Full-Adaptive Snakes introduce spatially varying contour parameters, RL-Adaptive Snake further incorporates temporal parameter scheduling through reinforcement learning, Multi-Scale Snake employs coarse-to-fine optimization, and Combined Snake integrates adaptive weighting with multi-scale optimization.

4.5 Implementation Details

All methods were implemented in Python using NumPy and SciPy.

For Adaptive Snake, the bending stiffness parameter was spatially adapted according to local image gradient magnitude. Full-Adaptive Snake further introduced spatially varying damping coefficients.

For RL-Adaptive Snake, a Deep Q-Network (DQN) controller was trained to select optimization parameters every 50 snake iterations. The action space consisted of five predefined (γ, σ) pairs:

$$(1,2),(2,3),(5,5),(10,6),(20,8). \tag{14}$$

The observation space included local gradient features, contour statistics, and optimization-state information. The reward was defined as the increase in IoU between consecutive decision stages.

All methods were initialized using identical contour configurations and evaluated under the same experimental conditions to ensure fair comparison.

5 Results

5.1 Adaptive Parameter Methods

We first evaluate the effectiveness of adaptive parameterization strategies, including Adaptive Snake, Full-Adaptive Snake, and RL-Adaptive Snake. These methods progressively introduce spatial and temporal adaptation mechanisms while maintaining the same underlying active contour framework.

5.1.1 Synthetic Benchmark Results

Figure 1 summarizes the quantitative results on the synthetic benchmark dataset.

For simple geometric shapes such as disks and ellipses, all methods achieve high segmentation accuracy with IoU values close to 1.0. The differences become more apparent on more challenging examples containing noise, weak boundaries, and non-convex structures.

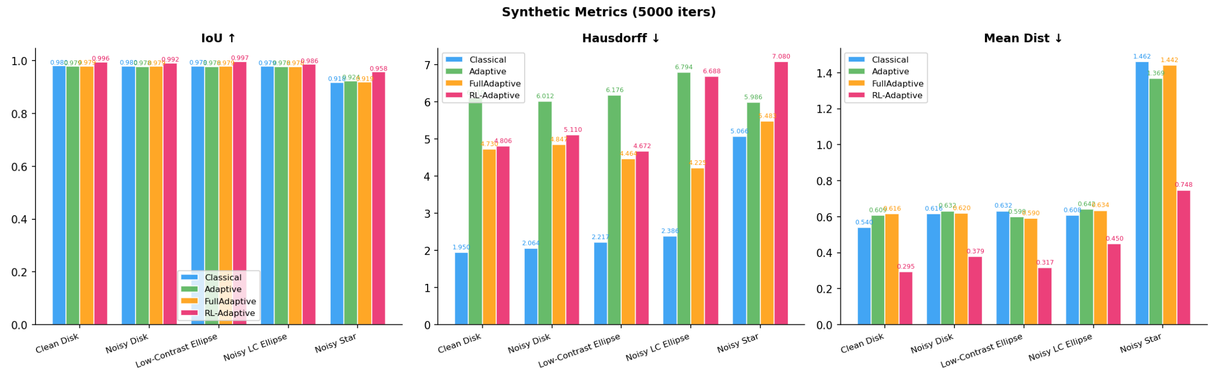

Figure 1: Quantitative comparison on the synthetic benchmark dataset.

Among all adaptive variants, RL-Adaptive Snake consistently achieves the highest IoU values across the benchmark. The largest improvement is observed on the noisy star example, where RL-Adaptive increases the IoU from 0.918 for the Classical Snake to 0.958. In addition, RL-Adaptive substantially reduces the mean boundary distance from 1.462 to 0.748.

Adaptive Snake and Full-Adaptive Snake improve contour flexibility through spatially varying parameters, allowing the contour to better accommodate local image structures. However, their optimization behavior remains fixed throughout contour evolution. By contrast, RL-Adaptive dynamically adjusts optimization parameters during the optimization process, enabling more aggressive exploration in earlier stages and more conservative refinement near convergence.

These results suggest that temporal adaptation provides additional benefits beyond spatial adaptation alone, particularly for objects with complex geometry and noisy boundaries.

5.1.2 Real-World Dataset Results

We further evaluate the adaptive parameter methods on the FLUO cell segmentation dataset and a medical ultrasound dataset.

On the FLUO dataset (Figure 2), all methods achieve similar IoU values, with RL-Adaptive producing slightly lower average overlap than the other variants. However, RL-Adaptive obtains the lowest Hausdorff distance, indicating improved worst-case boundary localization despite comparable segmentation overlap.

One possible explanation is that fluorescence microscopy images contain relatively compact and homogeneous cell nuclei, for which small differences in contour size can noticeably affect IoU while having only a limited impact on boundary accuracy. As a result, RL-Adaptive may produce contours that align more closely with difficult local boundary regions while sacrificing a small amount of overall region overlap.

These results suggest that temporal parameter scheduling can improve contour alignment in challenging local regions even when improvements are not fully reflected by region-based metrics such as IoU.

On the ultrasound dataset (Figure 3), RL-Adaptive achieves the strongest overall performance across all evaluation metrics. The average IoU increases from 0.615 for the Classical Snake to 0.644, while both the Hausdorff distance and mean boundary distance are reduced.

Qualitative examples shown in Figure **??** further demonstrate that RL-Adaptive produces contours that more closely follow weak anatomical boundaries while avoiding excessive smoothing. In several challenging cases, the RL-based controller better balances contour flexibility and stability, resulting in improved boundary localization.

5.2 Multi-Scale Optimization

We evaluate the Multi-Scale Snake on the synthetic benchmark to assess the effect of coarse-to-fine contour refinement under controlled shape complexity, noise, and low-contrast conditions. Because the Multi-Scale Snake preserves the classical contour parameters and modifies only the optimization schedule, its results isolate the benefit of progressive multi-resolution refinement.

On simple geometric structures such as disks and ellipses, the Multi-Scale Snake achieves consistently strong segmentation accuracy, with IoU values of approximately 0.980 across both clean and noisy variants. These results indicate that coarse-

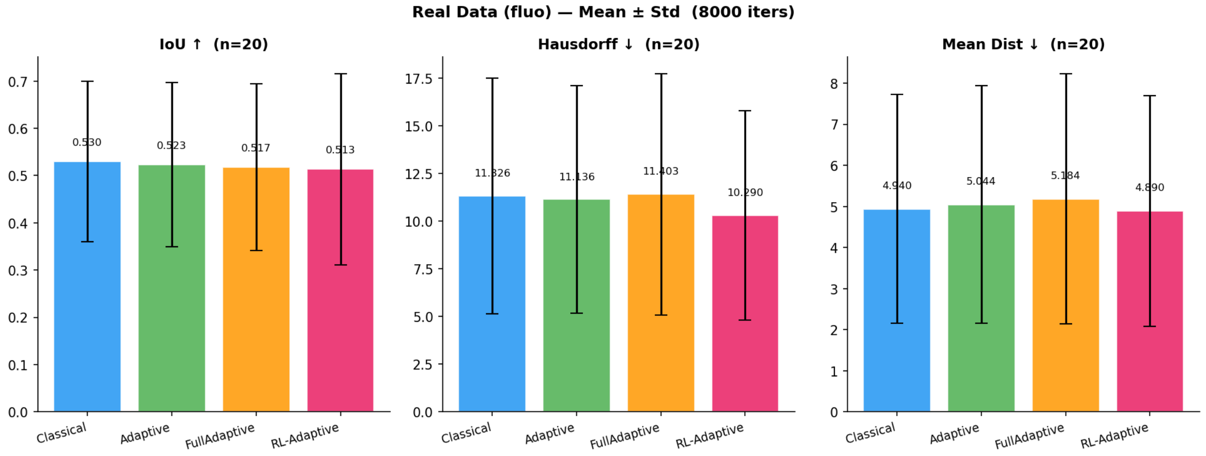

Figure 2: Quantitative evaluation on the FLUO cell segmentation dataset. Results are reported as mean $\pm$ standard deviation over 20 images.

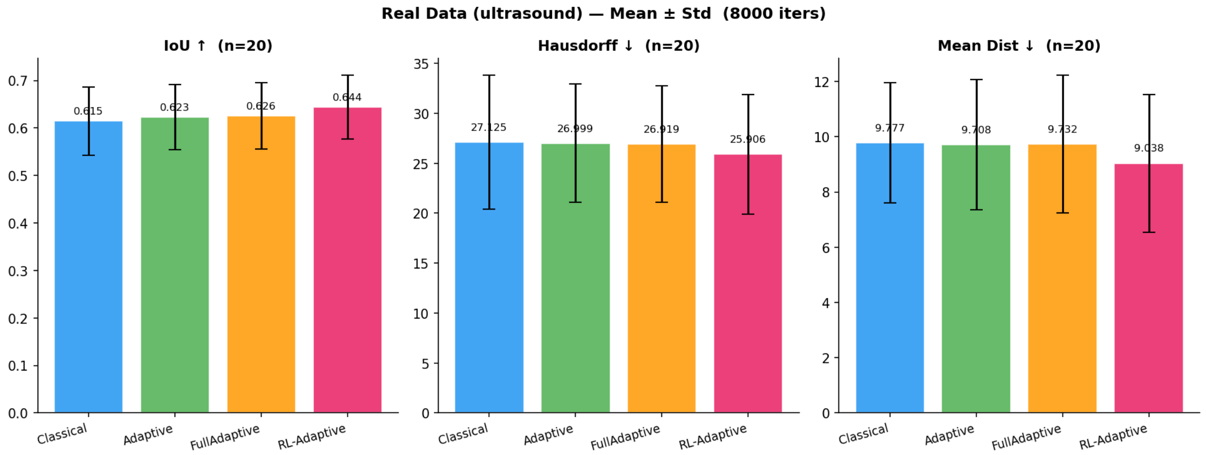

Figure 3: Quantitative evaluation on the ultrasound dataset. Results are reported as mean $\pm$ standard deviation over 20 images.

to-fine optimization preserves stable contour fitting under moderate perturbations while maintaining accurate region overlap.

The most challenging case remains the non-convex noisy star, where the contour must recover sharp corners and inward boundary regions. In this setting, the Multi-Scale Snake achieves an IoU of 0.926. Although this case remains more difficult than the smoother synthetic objects, the result shows that coarse-to-fine refinement provides a robust optimization baseline even without adaptive stiffness or temporal control.

Overall, the synthetic benchmark demonstrates that the Multi-Scale Snake improves optimization stability and remains effective under noisy and low-contrast conditions, while preserving the simplicity of the classical active contour formulation.

We additionally evaluated the Multi-Scale Snake on two real-world biomedical datasets. On the ultrasound dataset, the Multi-Scale Snake achieved a mean IoU of 0.438 ± 0.123, showing that coarse-to-fine refinement can provide moderate robustness even under weak boundaries and speckle noise. In contrast, performance on the FLUO dataset remained limited (0.070 ± 0.113 IoU), indicating that multi-scale optimization alone was insufficient for the fluorescence microscopy setting and remained sensitive to initialization and parameter choice.

Case	IoU	Hausdorff	MeanDist
Clean Disk	0.980	4.06	0.56
Noisy Disk	0.980	3.91	0.59
Low-Contrast Ellipse	0.980	4.33	0.58
Noisy LC Ellipse	0.979	3.16	0.60
Noisy Star	0.926	9.97	0.98

Table 1: Quantitative evaluation metrics of the Multi-Scale Snake on the synthetic benchmark dataset.

Metrics (mean $\pm$ std)	FLUO	Ultrasound
IoU	0.070 ± 0.113	0.438 ± 0.123
Hausdorff	46.8 ± 34.9	61.1 ± 27.8
MeanDist	21.43 ± 25.54	23.27 ± 11.50

Table 2: Quantitative evaluation of the Multi-Scale Snake on real-world datasets.

5.3 Combined Framework

The Combined Snake achieved superior accuracy across both synthetic benchmarks and complex real-world datasets by blending spatial shape flexibility with multi-scale stability.

5.3.1 Synthetic Benchmark Results

For less complex geometric shapes, such as the standard disks and ellipses, all evaluated active contour variants perform exceptionally well, yielding near-optimal segmentation scores with IoU values very close to 1.0 (typically around 0.980). When the target object structures are not complex, the spatial and temporal adaptation strategies produce highly similar and robust results across the board.

The distinct advantages of our combined framework become evident on more challenging test cases with high geometric complexity, such as the non-convex `noisy_star`. On this intricate shape, the Combined Snake successfully reaches the deep inward valleys of the star, scoring an IoU of 0.924 and lowering the maximum error boundary. Near the true boundaries, the local stiffness parameter dynamically drops to $\beta_{\min} = 0.005$, allowing the snake to flex into sharp corners while the multi-scale structure protects the global shape from background noise.

Case	IoU	Hausdorff	MeanDist
Clean Disk	0.980	4.06	0.56
Noisy Disk	0.980	3.91	0.59
Low-Contrast Ellipse	0.980	4.33	0.58
Noisy LC Ellipse	0.979	3.16	0.60
Noisy Star	0.924	10.19	1.04

Table 3: Quantitative evaluation metrics of the Combined Snake module on the synthetic benchmark dataset.

5.3.2 Real-World Dataset Results

We evaluate the Combined Snake framework on real-world biomedical images by applying modality-aware parameter routing.

- **Fluo-HeLa Dataset:** Fluorescence cell images feature dense groups of compact structures. By implementing localized Region of Interest (ROI) isolation masks and a micro-smoothing filter ($\sigma = 1.5$), the Combined Snake prevented cross-contamination from neighboring cells. This approach yielded a high accuracy of $\text{IoU} = 0.736 \pm 0.186$ and a very tight average localization error ($\text{MeanDist} = 1.47 \pm 0.74$ px).

- **Ultrasound Dataset:** Ultrasound scans suffer from severe speckle artifacts and missing acoustic boundaries. By deploying an expanded initialization radius ($\text{scale} = 1.3$) to start from a clean safety zone and utilizing a heavy macro-smoothing filter ($\sigma = 4.5$), the model ignored the surrounding muscle fiber texture noise. This configuration allowed it to capture the true outline of the nerve bundle, achieving a stable baseline mean score of $\text{IoU} = 0.631 \pm 0.088$.

Metrics (mean $\pm$ std)	Fluo-HeLa	Ultrasound
IoU	0.736 ± 0.186	0.631 ± 0.088
Hausdorff	3.7 ± 2.1	29.0 ± 10.5
MeanDist	1.47 ± 0.74	10.46 ± 4.03

Table 4: Quantitative evaluation metrics of the Combined Snake module across real-world biomedical imaging datasets ($n = 20$).

5.4 Discussion

The experimental results reveal that different adaptation strategies are beneficial under different imaging conditions. Temporal adaptation through RL-based parameter scheduling consistently improves performance on synthetic and ultrasound datasets, where optimization requirements vary throughout contour evolution. In these settings, dynamically balancing aggressive exploration and conservative refinement leads to more accurate boundary localization.

In contrast, the Combined framework achieves the strongest performance on the FLUO dataset. This suggests that fluorescence microscopy images benefit more from improved initialization and coarse-to-fine refinement than from temporal parameter scheduling. The relatively consistent object structures in this dataset may reduce the need for dynamic optimization behavior during contour evolution.

Overall, no single adaptation strategy is universally optimal. Temporal adaptation appears particularly effective for weak-boundary and noisy images, while multiscale optimization provides advantages in datasets with more consistent object appearance. These findings highlight the complementary roles of spatial adaptation, temporal adaptation, and multiscale optimization in robust active contour segmentation.

6 Conclusion and Future Work

We presented a comparative study of spatial, temporal, and multiscale adaptation for active contour segmentation. Our experiments show that reinforcement-learning-based parameter scheduling improves contour fitting on complex synthetic shapes and noisy ultrasound images, but does not consistently outperform fixed parameters on fluorescence microscopy data. Instead, the Combined framework, which integrates spatially adaptive stiffness with coarse-to-fine optimization, performs significantly better on the FLUO dataset, where cell boundaries are relatively well-defined and object structures are compact and consistent across frames, achieving the highest IoU and lowest Hausdorff distance among all evaluated methods. These results suggest that temporal adaptation is beneficial when optimization requirements change during contour evolution, while coarse-to-fine initialization provides greater benefit in datasets with more consistent object structure. Overall, no single strategy is universally optimal, and the effectiveness of each approach remains domain dependent.

Future work will investigate continuous parameter control, training on more diverse real-world images, and explore the potential of combining the RL-adaptive and current combined framework. Incorporating automatic initialization and convergence detection may further improve the robustness and efficiency of adaptive active contours.

References

[1] Michael Kass, Andrew Witkin, and Demetri Terzopoulos. Snakes: Active Contour Models. *International Journal of Computer Vision*, 1(4):321–331, 1988.

[2] Volodymyr Mnih et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, 2015.

[3] Ulman, V. et al. An objective comparison of cell-tracking algorithms. *Nature Methods*, 14(12):1141–1152, 2017.

[4] Kaggle. Ultrasound Nerve Segmentation. https://www.kaggle.com/c/ultrasound-nerve-segmentation, 2016.

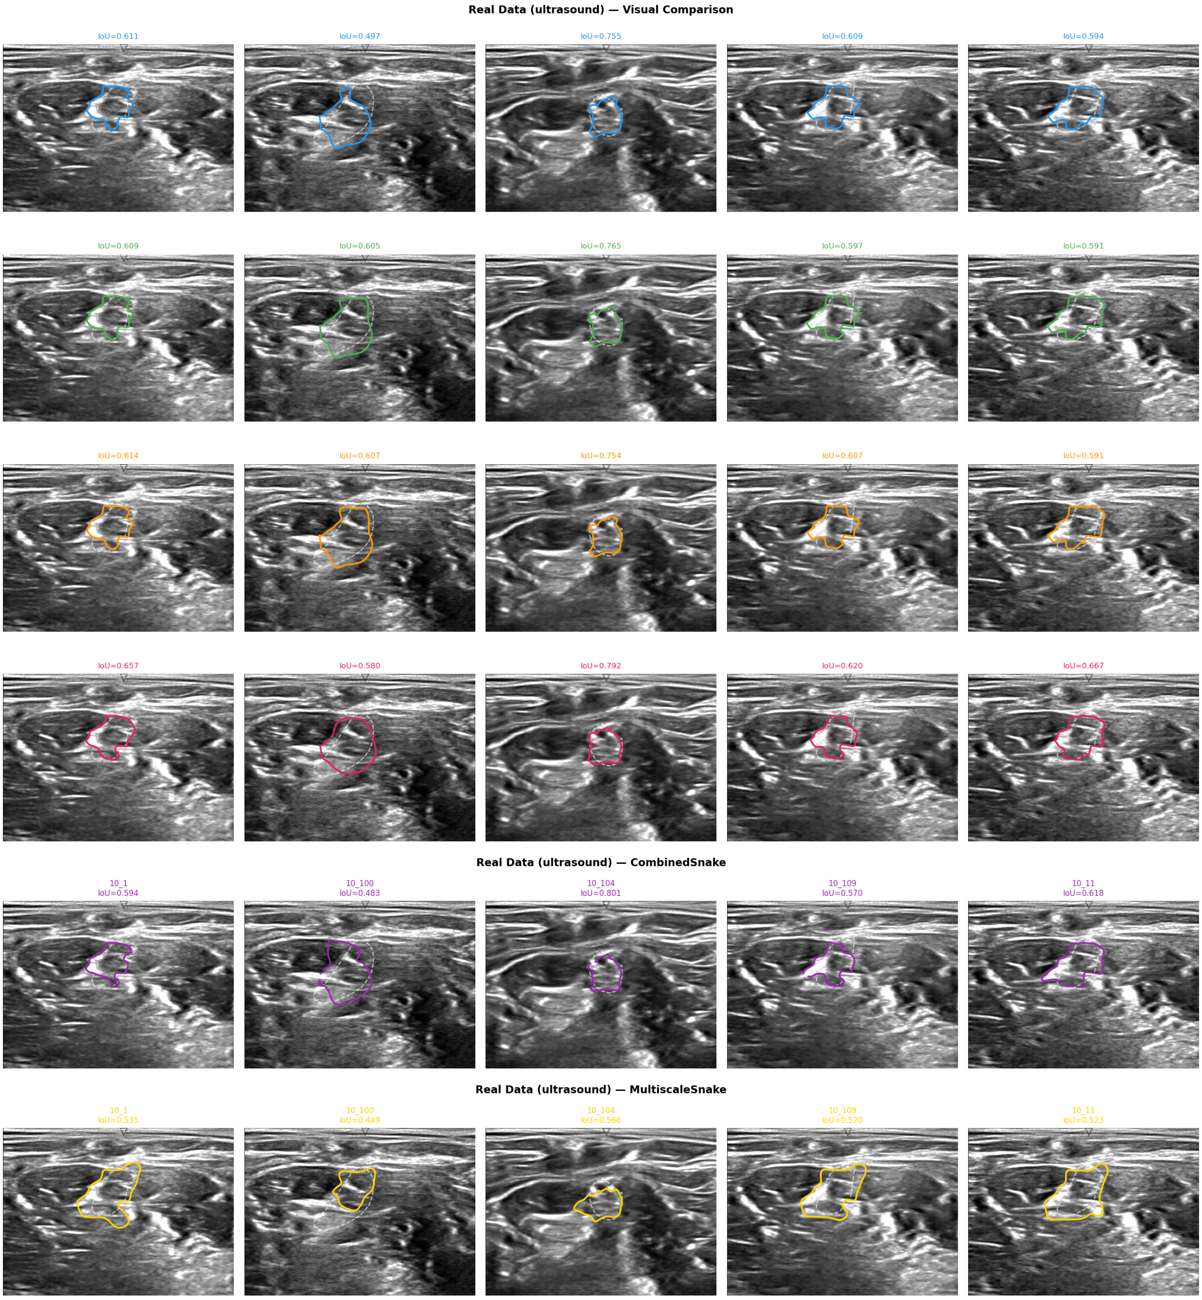

Figure 4: Representative contour fitting results across different datasets.